

Assignment 05: LC-3 OS System Call Traps/Exceptions and Memory Mapped I/O

CSci 450: Computer Architecture

Objectives

In this assignment you will continue our work on the LC-3 ISA simulator that we started previously, and continue practicing C. Previously we implemented most of the 16-bit LC-3 microarchitecture. In this assignment we will add onto the LC-3 simulator the mechanisms to support some of the OS-level features defined for the LC-3 computer architecture.

In particular, in this assignment you will get an example of one way that hooks for invoking system calls can be created for an OS to run on time of the ISA/microarchitecture. In this assignment you will implement the trap call/return mechanism. This will require us to add in priority and privilege level mechanisms to the simulated computer architecture. In essence, you will implement a supervisor/user mode, and the mechanisms that switch between these when a trap, exception or interrupt happens.

We will also implement some basic memory mapped I/O in this assignment, so that we can provide keyboard input and display character output on a terminal from the LC-3 simulator. You will use the memory mapped I/O to create and test basic service routines to get and display character I/O from the trap service routines.

Questions

- How are privilege and priority levels defined and implemented in a simple computer architecture?
- What is the relationship between a trap, an exception and an interrupt?
- How are system calls implemented by computer architectures to allow for an operating system to implement additional levels of abstraction for a computer?
- What are typical mechanisms to implement memory mapped I/O operations in a computing system?
- How are access to I/O devices provided as higher level abstractions by operating system service routines?

Objectives

- Become more familiar with the LC-3 ISA definition.
- See an implementation of operating system level abstractions in a simple LC-3 ISA simulation.
- Understand how privilege works in the OS level and its use and relationship to os system calls.
- Understand how basic memory mapped I/O functions in a simple computer architecture.
- Learn about system level C functions useful in implementing procedures at this level of a machine description.

Description

In this assignment you will be extending the previous implementation of the LC-3 microarchitecture simulation to add in several OS-level computer architecture concepts and mechanisms. You have been given a working implementation of our previous microarchitecture assignment as a starting point, though feel free to use or compare with your previous assignment implementation while adding in the OS-level mechanisms for this assignment.

Overview and Setup

For this assignment you will be adding in additional functions and declarations for the `lc3vm.cpp` simulation. As usual before starting the assignment tasks proper, you should make sure that you have completed the following setup steps:

1. Accept the assignment and copy the assignment repository on GitHub using the provided assignment invitation link for 'Assignment 05: LC-3 OS Level Mechanisms' for our current class semester and section.
2. Clone the repository using the SSH URL to your host file system in VSCode. Open up this folder in a Development Container to access and use the build system and development tools.
3. Confirm that the project builds and runs, though no tests will be defined or run initially. If the project does not build on the first checkout, please inform the instructor. Confirm that your C/C++ Intellisense extension is working, and that your code is being formatted according to class style standards when files are saved.
4. You should create the issue for Task 1 and/or for all tasks for the assignment now before beginning the first task. On your GitHub account, go to **Issues**, and create them from the issue templates for the assignment. Also make sure you are linking each issue you create with the **Feedback** pull request for the assignment.

Assignment Tasks

The starting state for this assignment has implementations of the simulated LC-3 microarchitecture functions to implement the basic fetch-decode-execute cycle for the LC-3. The basic outline of the simulator functions like this at this point:

- We load 1 LC-3 machine language (binary) program into the (simulated) main memory.
- In the RPC register, we keep the program counter, the next instructions to fetch, decode and execute.
- The instruction is fetched and the **Operation Code** (first 4 bits) is obtained from this instruction. Based on that, we decode the rest of the address and control bits from the fetched instruction.
- We execute the method associated with the given instruction, using a function lookup table.
- We increment the RPC for normal sequential execution and continue with the next fetch-decode-execute cycle.

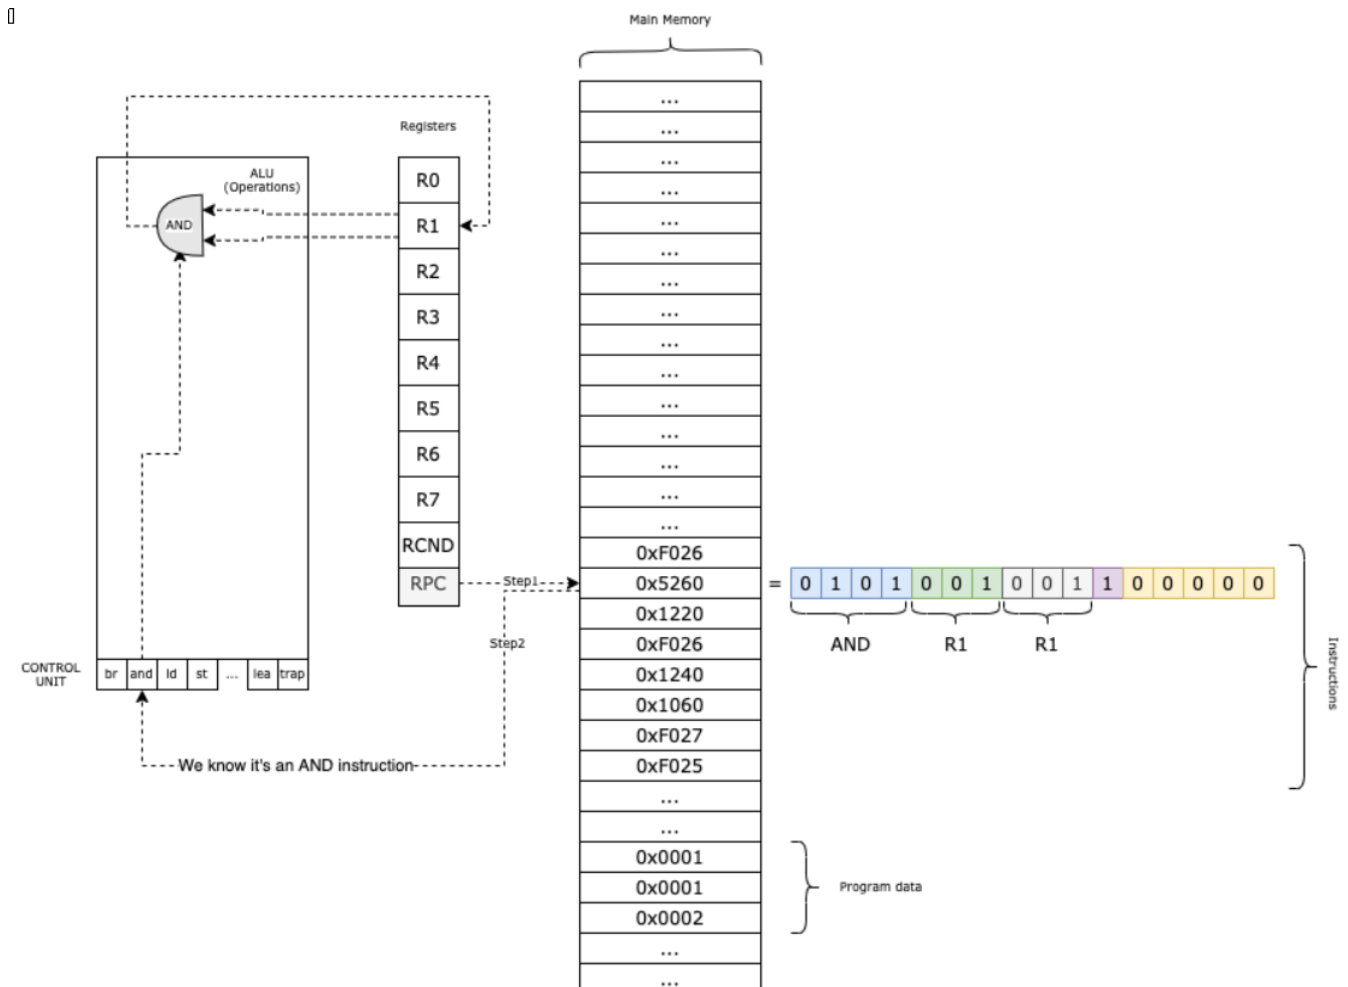


Figure 1: Example LC-3 Fetch-Decode-Execute

Task 1: Processor Status Register PSR Privilege and Priority

Task 1 Part 1: PSR Privilege Mode Bit

For task 1 we are going to add in additional meaning to the Processor Status Register (PSR). Previously we had only made use of the low 3 bits of the PSR to hold the N,Z,P condition codes. In order to implement traps, exceptions and interrupts, we need to add in a supervisor/user privilege mode distinction to the computer architecture, and a notion of priority level for interrupts. The bits we will use in the PSR look like the following:




Figure 2: Processor Status Register (PSR)

Recall that we already define `reg[PSR]` to hold the current process status register bits, and there is a macro named `RCND` which masks the low 3 bits so that we can test if the N, Z or P flags are set easily.

As usual start by defining `task1_1`. We have broken up the first task into 2 parts. For the first part you need to implement a few methods to set and test the privilege mode bit of the PSR. To simplify the implementation, it is better to define all of the following as regular C functions, rather than using C macros.

When bit 15 of the PSR is a 0 the machine is in supervisor mode (the more privileged mode used by the OS). When bit 15 is 1 we are in user mode.

You should see from the tests for task 1 part 1 that we start by setting all of the bits of the PSR to 0, thus we are initially in supervisor mode.

You need to write 3 functions to get the task 1 part 1 tests to pass. You need to declare each function in the `lc3vm.h` header file, and put the implementation into `lc3vm.c`.

1. The `is_user_mode()` function should return a `bool` result. It needs to test bit 15 and return `true` if the bit is 1 indicating running in user mode, and return `false` if it is 0 indicating currently in supervisor mode. You need to use masks here as the N/Z/P and priority level bits may not always all be 0.
2. `user_mode()` is a `void` function. It should set the bit 15 of the PSR to be 1, leaving all of the other bits unchanged. The tests start by switching into user mode by calling this function.
3. `supervisor_mode()` is also a `void` function that puts the machine into supervisor mode by setting bit 15 to 0.

Once you implement these 3 functions they should successfully pass the tests given for part 1 of task 1.

Task 1 Part 2: PSR Priority Level Bits

Bits 10-8 of the PSR keep track of the priority level that the machine is currently running at. Priority levels are mainly used for the interrupt handling mechanisms in computer architectures. Since 3 bits are used for the priority level, this means 8 level of priority from 000 (0) to 111 (7) are possible in the LC-3 machine.

NOTE: Read Patt and Patel Chapter 9.1 for details on the privilege and priority level mechanisms.

Define the `task1_2` tests and implement the following

1. `priority()` returns the current priority level set in the PSR priority bits. This should return an `uint16_t` type.
2. `set_priority()` is used to set the priority bits. Again pass in a `uint16_t` type for the parameter, though we only use the least significant 3 bits of this parameter value.

These function may be a bit trickier than the previous ones, especially `set_priority()`, since you are working with 3 bits for the priority level. For `set_priority()` one approach is to first 0 out the PSR[10:8] bits, so that you can then shift up the passed in parameter and or (`|`) it with the PSR in order to set the needed priority level bits.

The tests for `task1` test not only that you set or get the appropriate bits, but also that you are not modifying bits that should not be changed in the PSR when you change the privilege or priority levels.

Once you have all of the tests passing for parts 1 and 2 of task 1, create a commit and push it to your GitHub classroom repository.

Task 2: Stack push/pop for trap/exception mechanism

Note: Read Patt & Patel section 9.1.2 on the LC-3 Organization of Memory and appendix A.2 and A.3 for descriptions of the trap, interrupt and exception processing mechanisms.

With the addition of a user/supervisor privilege mode mechanism, we can begin to add in other needed mechanisms so that an operating system can manage a computing system and allow user programs to run on the system and make higher level services available for use by user programs.

Recall that the LC-3 computer uses a 16-bit memory address space. The layout of LC-3 memory looks like the following:

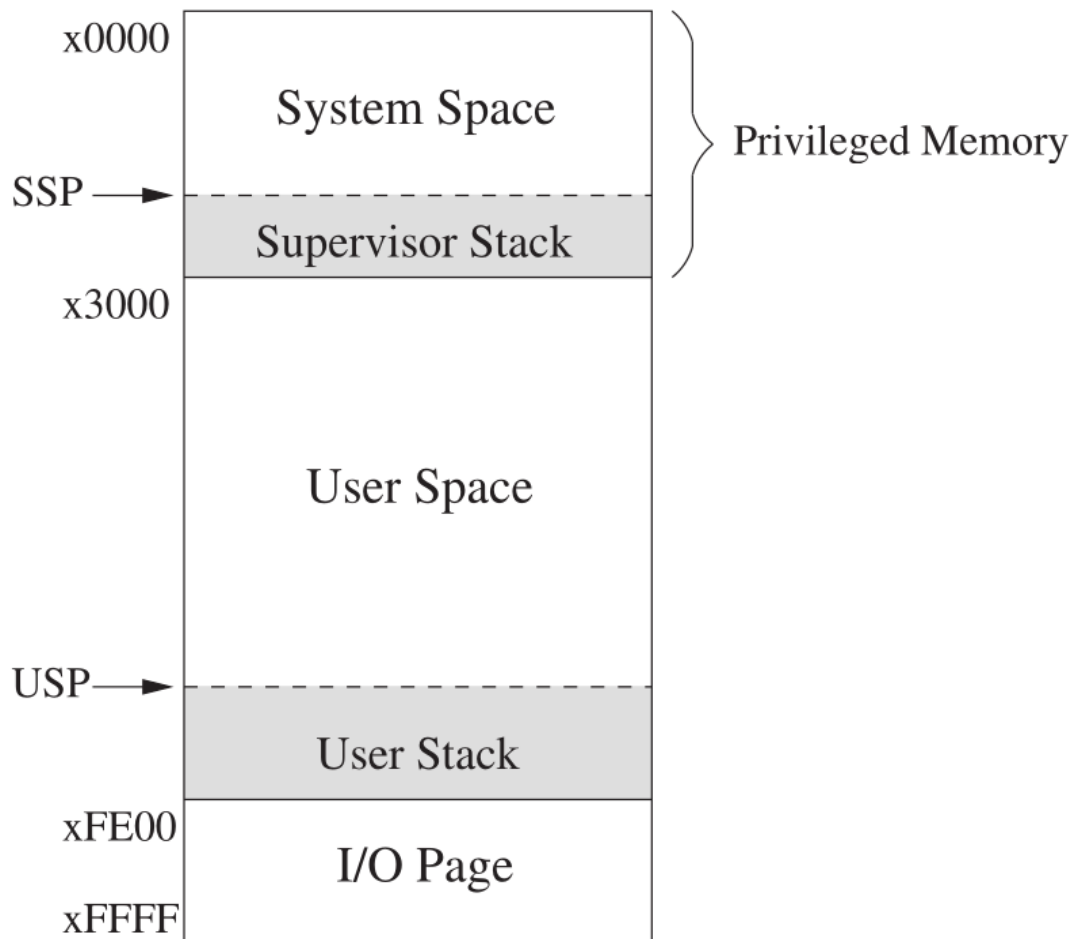


Figure 3: LC-3 Regions of Memory

The space from `0x3000` to `0xFE00` is user space. Programs that run in regular user mode are loaded into user space. System space is defined from `0x000` to `0x2FFF` in memory. This is where the operating system loads its programs. In order to run code that is in the system space privileged memory, the machine must be in supervisor mode.

Notice that at the bottom of system and user space stacks are defined. We will need to add in two additional registers in a later task to our LC-3 register declarations, the **SSP** and **USP** which will hold the top of the System Stack Pointer (SSP) and the User Stack Pointer (USP) respectively. Initially USP is initialized to `xFE00` and SSP is initialized to `0x3000`.

The stacks are used for jumping into and returning from subroutines (JSR, JSRR/RET) and for invoking traps, exceptions and interrupts.

By convention in the LC-3 machine, R6 is treated specially. This register holds the memory address of the top of the current stack being used. The SSP and USP we mentioned are used when a mode switch from user to supervisor mode or back occurs. For example when we switch from user mode into supervisor mode, R6 should have the stack pointer for the current user stack. We save R6 in the USP and then load the SSP into R6 when we switch to supervisor mode, so that now we are pushing and popping onto the supervisor stack. The reverse occurs when we mode switch from supervisor mode back to user mode.

For this task, we first want to implement convenience methods that manipulate R6 for pushing and popping values on and off whichever stack it is pointing to.

In the figure showing the LC-3 memory layout, the smallest address 0x0000 is at the top and goes down to the largest address 0xFFFF at the bottom. By convention, R6 holds the memory address of the top of the current stack, and when we push values onto the stack, the address needs to be subtracted by 1 when growing the stack, and when we pop a value the stack needs to have 1 added to shrink the stack.

Define the `task2` tests and implement the push and pop microarchitecture methods that will be used to manipulate the supervisor and user stacks through the R6 register. Do the following:

1. Implement `push()`, which will take an `uint16_t` value as its input parameter (and it is a `void` function). This method assumes that R6 is pointing to the top of the current stack in use. Subtract 1 from the stack address in R6 to grow the stack, and then place the parameter that was pushed into memory that is now pointed to as the top of the stack by R6. You are required to reuse the `mem_write()` function here, as we later need to add in checks that we are not accessing privileged memory when in user mode.
2. You also need to implement `pop()` which takes no parameters and is also a `void` function. This function should pop off the value from the top of the stack. To do this you simply need to add 1 to the current value of the R6 register holding the current stack pointer.

Task 3: Master Control Register (MCR) and halting execution

Note: Read Patt & Patel section 9.3 on OS Service Routines, in particular we first start here with section 9.3.7 to add in the RUN latch of the MCR

Recall that in our first implementation of the LC-3 virtual machine we had simply defined a global `running` boolean variable that the halt trap service routine set to `false` in order to cleanly stop the fetch-decode-execute cycle. Years ago many ISAs had a separate HALT instruction, but this was very rarely executed so seemed wasteful to use an entire opcode for it. In modern machines a RUN latch, or clock enable bit, is memory mapped to an address in the computer, and any system service routine can halt the execution of the computer by clearing a particular bit that controls the RUN latch. In LC-3 the RUN latch bit is bit [15] of the Master Control Register, which is memory mapped to location 0xFFFFE.

Note: See Patt & Patel Table A.1 for the device register memory mapping assignments.

We are going to simulate the memory mapped MCR and manipulating its bit [15] to control the clock enable/disable of our machine.

Note: You should look at the changes made to the definition of the `reg[]` array. Notice it is now a pointer and we declare it in `lc3vm.c` like this:

```
// memory map the registers starting at address 0xFFFF0, maps PSR to 0xFFFFC and MCR to 0xFFFFE
// uint16_t reg[RCNT] = {0};
uint16_t* reg = &mem[0xFFFF0];
```

So really all of the registers are mapped into memory addresses starting at address 0xFFFF0. This includes that the PSR processor status register is actually mapped to address 0xFFFFC and that the MCR machine control register we are about to use is mapped to address 0xFFFFE as specified in the LC-3 architecture.

Implement the following functions to test and set the MCR clock enable bit:

1. `enable_clock()` is a `void` function that takes no parameters. This function should simply set the MCR bit [15] to 1 in order to enable the system run latch. You are required to use the `reg[]` mapping to access MCR from this function.
2. `disable_clock()` is also a `void` function with no parameters, and it should do the opposite and set the MCR bit [15] to 0.
3. `is_running()` is a function that returns a `bool` result by testing the MCR run latch bit [15]. It should return `true` if the clock is currently enabled and thus the system is running, and `false` if not. We will later be using this function to control the fetch-decode-execute loop termination.

Once you are passing all of the task 3 tests and are satisfied with your implementation, make sure you create a commit for task3 and push your work to your GitHub classroom repository.

Task 4: `trap()` and `rti()` opcode implementations

Note: Read Patt & Patel section 9.3 for information about OS system service routines in LC-3 using traps. But also read the appendix A.2 details of the `trap` and `rti` instructions for details on how these need to be implemented.

Task 4 Part 1: (re)Implement `trap()` microarchitecture opcode

The `trap` instruction in LC-3 is how we allow user programs to invoke system / service calls implemented by the operating system. In particular a `trap` will first cause the computer to switch to supervisor mode if we are currently in user mode, and it will then be able to run a trap service routine provided by the operating system in the privileged mode.

If you look at the code given for this assignment, the implementation we had given for traps previously has been removed, and the `trap()` and `rti()` functions are currently empty functions. We need to reimplement traps using the new privilege level modes and functions.

The pseudocode given in A.2 for the `trap` instruction is as follows:

```
TEMP = PSR;
if (PSR[15] == 1) // e.g. if currently in user mode
    USP = R6 and R6 = SSP
    PSR[15] = 0 // change to supervisor mode now
push PC
push TEMP // push PC and the original PSR on stack
PC = mem[ZEXT(trapvect8)];
```

In English, when a `trap` occurs, if we are currently running in user mode, we save the current value in `R6` which has the current user stack pointer, and we restore the system stack pointer into `R6`. Once this is done we set the MCR bit [15] to 0 so we are now running code as supervisor. Notice that the switch of the stack pointers only needs to be done when we mode switch from user into supervisor mode. Whether we switch mode or were already in supervisor mode, we need to push the `PSR` (before we possibly modified the `PSR` privilege bit) and the current `RPC` onto our stack, which since we are guaranteed to now be in supervisor mode should be pushed onto the supervisor stack.

Finally what the `trap` really is doing is calling a trap procedure. We need to save the current `PSR` and `RPC` on the stack so that if/when an `rti` instruction executes to return from the trap/interrupt, we can restore the processor status and the `PC` to continue executing from where the trap was invoked. So the last statement modifies the `PC` to change executing to the invoked trap service routine. The low 8 bits [7-0] of the `trap` instruction are really an index into the trap service routine table. The trap service routine table is loaded into memory addresses `0x0000 - 0x00FF`. So the LC-3 architecture can support up to 256 trap service routines as operating system service calls.

The trap service table in addresses `0x0000 - 0x00FF` is really only a lookup table of addresses to the location of the service routine. The operating system will usually load the service calls that it defines into protected memory, and set up the trap service routine mechanism so they can be invoked using the `trap` instruction with an appropriate `trapvect8` code.

Define the `task4_1` tests and implement the `trap()` function to pass the given tests. You need to implement the pseudo code given above for the `trap` instruction. Recall that the instruction is passed in as a parameter to this function (the declaration of `trap()` and the empty function should already be in the file for you for this task).

Some hints about the implementation:

- You are required to reuse the `is_user_mode()` and `supervisor_mode()` functions to test and switch from user mode to supervisor mode.
- Likewise you need to use the `push()` method to push on the PSR and the PC to the system stack.
- You are required to push the PC first then the original PSR, so that the PSR ends up on top of stack, and the return PC below it.
- As shown in the pseudocode, you need to save the PSR before switching to supervisor mode, the PSR that is pushed on needs to be what it was when the function is called.
- There is a `TRP` macro defined in `lc3vm.h` that you can and should use to extract the low 8 trap vector bits from the instruction.
- You are required to use `read_memory()` here to look up the trap service routine address in memory using the `trapvect8` bits of the instruction.

You should notice in the tests that the `ld_img()` function is used to load an assembled LC-3 image into memory before the tests. This contains the address of the trap service routines used in this set of tests. You might want to look at the `progs/trap-vector-table.asm` file to get an idea of how this image code is created and assembled.

Complete the `trap()` implementation to pass the given tests before moving on to part 2 to implement the corresponding return from interrupt/trap instruction.

Task 4 Part 2: (re)Implement `rti()` microarchitecture opcode

The `rti()` instruction is the logical inverse of invoking a service call using `trap`. Note that `rti` is a mnemonic for “return from trap or interrupt. The interrupt service routines have basically the same mechanism in LC-3, and we can return from an interrupt or trap service routine using this opcode.

The pseudocode given for `rti` in appendix A.2 of the Patt and Patel book is:

```
if (PSR[15] == 1) // if we are in user mode you can invoke rti
    Initiate a privilege mode exception

PSR = mem[R6]    // R6 is the SSP, the PSR is restored
pop
PC = mem[R6]     // R6 is the SSP, PC is restored
pop

if (PSR[15] == 1) // we did a mode switch back to user mode
    SSP = R6 and R6 = USP
```

We will later in the assignment add in exceptions, so you don’t have to handle generating an exception if `rti` is called while in user mode.

Otherwise the `rti` instruction is relatively simple, you pop off and restore the PSR and RPC from the system stack. When you restore the PSR you may be restoring back to user mode privilege status, indicating you are returning back to userland from a user invocation of a trap. In that case you need to save the system stack pointer and then restore back the saved user stack pointer to R6.

- You are again required to use the `pop()` methods to manipulate the system stack.
- And you should be using `mem_read()` to read the PSR and RPC values when restoring them from the stack.

Implement the `rti()` function to pass the given unit tests. Once you are passing your tests for `trap` and `rti` instructions, create a commit and push your work back to your GitHub classroom repository.

Task 5: halt trap Clock Disable

So far we have added in the MCR machine control register bit to enable and disable the simulated system clock, and a working `trap` invocation and return mechanism. We are ready to begin adding in functionality for the LC-3 simulator and prepare to implement memory mapped I/O. Before we move on, let’s modify the existing `start()` function, that implements the basic fetch-decode-execute cycle, to use the new run latch control bit.

Enable the `task5` tests, and modify the `start()` method to use the new `is_running()` method that is testing the MCR clock enable bit. The `task5` tests load the trap vector table again and a simple `halt` trap service routine. Examine the file named `progs/task5-halt-trap.lc3`. When invoked this service routine will set the run latch bit to 0. So if you modify the main `start()` loop to be testing this bit, it should successfully terminate the call to `start()` in the test for this task. If you don't have this working properly, your tests might enter into an infinite loop.

- You are required to reuse your `is_running()` method in the `start()` function to test the clock run latch.

Once you are able to pass the tests for this task, create a commit and push your work to your GitHub classroom repository.

Task 6: KBSR and DSR I/O Status Updates on Memory Access

Note: consult Pat & Patel chapter 9.2 for more details about I/O in LC-3 and how memory mapped I/O functions. Also Figure A.1 again for memory layout and Table A.1 for the device register assignments.

We are now going to add in memory mapped I/O functionality for a (simulated) keyboard input device and a (simulated) output display device. As we saw earlier in the figure of the memory layout for LC-3, addresses in the range from `0xFE00` thorough `0xFFFF` are reserved for device register mapping addresses. We have been using registers and the PSR and MCR that have been mapped into this address space.

If you look in `lc3vm.h` and `lc3vm.c` there is also defined an `iomap[]` array that maps to the I/O device range of memory. The Table A.1 for device register assignments is shown here:

Table A.1 Device Register Assignments		
Address	I/O Register Name	I/O Register Function
<code>xFE00</code>	Keyboard status register (KBSR)	The ready bit (bit [15]) indicates if the keyboard has received a new character.
<code>xFE02</code>	Keyboard data register (KBDR)	Bits [7:0] contain the last character typed on the keyboard.
<code>xFE04</code>	Display status register (DSR)	The ready bit (bit [15]) indicates if the display device is ready to receive another character to print on the screen.
<code>xFE06</code>	Display data register (DDR)	A character written in bits [7:0] will be displayed on the screen.
<code>xFFFF</code>	Processor Status Register (PSR)	Contains privilege mode, priority level and condition codes of the currently executing process.
<code>xFFFE</code>	Machine control register (MCR)	Bit [15] is the clock enable bit. When cleared, instruction processing stops.

Figure 4: Device Register Assignments

You can access the device mappings using the array and defined enumerated type, for example to get the keyboard status register do `iomap[KBSR]`. There are also some defined constants of the actual memory address of the device registers, for example `KBSR_ADDR` is the keyboard status register address `0xFE00`.

We will start adding support for the I/O devices but hardcoding the needed updates of the memory mapped status registers on a read or a write of the device.

Task 6 Part 1: Update KBSR when KBDR is read

The way that memory mapped polling works is that the high bit [15] in the status register is used to signal when the device is ready. For the Keyboard device, the KBSR bit [15] should be set to 1 when a character is read and waiting in the KBDR. This will happen later as part of the fetch-execute-decode cycle where we will check for input. But the KBSR bit [15] needs to be cleared to 0 whenever the KBDR memory address is read by some program.

So add the following logic to the `mem_read()` function:

```
if address is KBDR_ADDR
    clear KBSR[15]
```

Of course you should do this before returning the value that is read from memory. The following are requirements

- You need to test the address against the `KBDR_ADDR` to see if the bit should be cleared
- But you should be using the `iomap[KBSR]` when manipulating and clearing the keyboard status bit.

- Make sure you do this before reading and returning the asked for value.

Once you can successfully clear the KBSR bit on a read of the KBDR, continue on to the second part of the task.

Task 6 Part 2: Update DSR when DDR is written

Likewise for the display device, we will hardcode in the basic fetch-decode-execute cycle to see if the DSR is 0, meaning that there is a character in DDR that should be displayed on the display device. The LC-3 supervisor programs will then display characters by writing a character into the DDR. So on a write to the DDR we need to again clear the DSR to 0. This signals that the display hardware needs to display the character, and once that is done it will set the DSR back to 1 in the fetch-decode-execute cycle.

So add the following logic to the `mem_write()` function:

```
if address is DDR_ADDR
    clear DSR[15]
```

Keep in mind the following

- As hinted at again you should be using the defined `DDR_ADDR` when testing if the display data register is receiving a new character.
- And also use `iomap[DSR]` to manipulate and clear the display status register bit.

Once your implementation is satisfactory and passing the tests for both parts, make a commit of your task 6 work and push it back to your GitHub classroom repository.

Task 7: Exception Handling

Note: See Pat & Patel section A.3.2 for details on exceptions in the LC-3 architecture and handling exceptions.

There are 3 exceptions defined in the LC-3 ISA: illegal opcode, privilege mode violation and access control violation (ACV).

The illegal opcode exception is generated if the unused reserved opcode (bits [15:12] = 1101) attempts to be executed.

A privilege mode violation occurs if an instruction that can only be used in supervisor (privileged) mode is attempted while running in regular user mode. For LC-3 the RTI instruction is the only one that is reserved completely for supervisor mode execution, to return from an interrupt, exception or trap service routine. The reason for this is that the RTI might cause a mode switch back to user mode, and in general userland code cannot run service routines that need to change back to user mode.

An access control violation (ACV) occurs if an attempt is made to read or write privileged memory space while the machine is in regular user mode. As indicated in previous drawings of the memory layout for the LC-3 architecture, all addresses from 0x0000 through 0x2FFF are privileged address. But also, not really indicated in the figures, but all memory mapped I/O and register addresses are also considered privileged memory, the address space from 0xFE00 up to 0xFFFF.

The handling of an exception works similarly to how a trap (and also an interrupt) is handled. One difference is that exceptions and interrupts are hardware device generated, while a trap is software generated by user programs. So the LC-3 microarchitectures will generate exceptions when needed, but our implementation in our simulation will be very similar to the `trap()` instruction function.

The exception supplies an eight-bit vector, like the trap vector. The exception vector table resides at addresses in memory starting at 0x0100, just above the trap vector table. The following three exceptions are hardcoded in the LC-3 specification:

Vector	Address	Exception
0x00	0x0100	privilege mode violation
0x01	0x0101	illegal opcode exception
0x02	0x0102	access control violation (ACV)

Task 7 Part 1: `except()` microarchitecture method

Start by copying the declaration and implementation of the `trap()` instruction function, and renaming the copy `except()`. The only modifications you need to make, besides changing the name and updating the function documentation, is that the exception vector needs to be looked at from an offset of 0x0100, so you need to add that to the exception vector passed in to this function to look up the location of the exception handling routine. You can reuse the `TRP` macro to extract the 8 bits, but simply add the needed offset to this vector.

Once you enable `task8_1` tests and add the `except()` method, it will load an exception table and test that your function is still correctly changing to supervisor mode and manipulating `R6` and the stack when an exception is invoked. Once the part 1 test for this task are passing, move on to the next parts to add in calls to generate the exceptions where needed.

Task 7 Part 2: Privilege Mode exception

When implementing the `rti()` instruction, there was a mention in the pseudocode that if we are currently in user mode that an exception should be thrown. Add in the test at the beginning of the `rti()` instruction function that will check if we are in user mode, and if so will call `except()` to generate a privilege mode violation (except vector 0x01).

NOTE: You can't continue on after calling the `except()` here, you need to return and not finish the rest of the `rti()` instruction on the privilege mode exception.

Enable the `task8_2` tests once you have added this to see if invoking RTI instruction while in user mode causes the exception.

Task 7 Part 3: Illegal opcode exception

Similarly if an attempt to execute the reserve opcode happens, this exception should be generated by the (simulated) hardware. Add in the appropriate call to the `res()` microarchitecture function to throw an illegal opcode exception. Then enable the `task7_3` tests and check that it is thrown if an attempt is made to execute the reserved opcode.

Task 7 Part 4: Access Control Violation (ACV)

Finally we will implement basic memory protection. We want to protect any read or write access to privileged memory when the machine is currently in user mode. So add in checks to both `mem_read()` and `mem_write()`. If the machine is currently in user mode, and the address is in the privileged address space (0x0000 - 0x2FFF system space or from 0xFE00 to 0xFFFF for privileged device register addresses). To simplify you can test if the address is either below 0x3000 or above 0xFDFD.

Add in these tests to `mem_read()` and `mem_write()`. In both cases the operation should not be performed, e.g. `mem_write()` should not cause a change in privileged memory to occur, and `mem_read()` should just return a 0x0 value. The access control violation ACV exception should be invoked for these cases and the function should return without writing or return 0x0 after the exception call. Note this should also happen before any test/updates of the device status registers occurs as well.

Enable the `task7_4` tests once you have implemented the checks for access control violations. Once all tests for task 7 are passing, make a commit of your task 7 work and push the work to your GitHub classroom repository.

System Tests: Testing and Running the LC-3 Simulator

For this assignment, the final 10 points are given for having all of the system tests passing successfully. There is one line of code that was commented out in the `start()` function, that calls `check_device_status()`. We commented this out when developing tasks 1-7 so that it would not interfere with the development of the tasks. We had also commented out calls to the `enable_clock()` and `user_mode()` functions in the `init()` initialization method, but you should have enabled those as described before performing the task 5 to enable and test the hat trap mechanisms.

But still it is easy to have some small mistakes in your implementations of tasks that, while they pass the unit tests, will not completely work when performing full system tests where we load an LC-3 program and run it. Uncomment the final call to the `check_device_status()` and try running the given system tests.

You can run the system tests by opening a terminal and doing a

```
$ make system-tests
```

Which will invoke a script that loads several test programs, executes them and then checks their results. If all of your system tests are passing, then your implementations are probably good and you can make a final commit of your code and push it to your GitHub classroom repository. If not all tests are passing, you will need to do some further investigation and debugging of your work.

Whether your tests pass or not, it is worthwhile to take a look at the code that is run for the system tests. For example, you can run the `prog-04` test like this:

```
$ ./lc3vm progs/prog-04.exe
```

The `prog-*.exe` files in the `progs` subdirectory contain linked together images. You should examine all of the `.lc3` assembly files in this subdirectory and understand what is being done here. The LC-3 assembly files define the `trap-vector-table.lc3` and the `exception-vector-table.lc3` which simply fill in hooks to the trap and exception system service calls provided by our primitive LC-3 architecture OS. You will find implementations of the traps in the `trap-*.lc3` files, and of the exceptions in the `exception-*.lc3` assembly files.

For example, if you tried running `prog-04.exe` by hand, examine the `prog-04.lc3` assembly file to see what it does. It simply invokes `trap 0x23` 3 times, in order to call the `trap-in.lc3` trap system service routine. This routine writes a prompt to the display device and reads in a single character. The `prog-04.lc3` ends by invoking `trap 0x25` to call the `trap-halt.lc3` service routine. All-in-all the `prog-04` is testing the IN trap service routine, which uses the keyboard I/O memory input mechanisms of our LC-3 machine. And all of the routines send output to the display device, testing the memory mapped output mechanisms. The `prog-01.lc3`, `prog-02.lc3` and `prog-03.lc3` programs simply cause each of the 3 different exceptions to happen, in order to test the exception handling mechanisms.

Please do run the programs by hand and make sure you understand how they work. If any of the tests are not passing once you uncomment the final code in the `start()` function, you will need to debug the execution of the LC-3 vm simulator to try and determine where the issue lies in your implementation.

Assignment Submission

For this class, the submission process is to correctly create pull request(s) with changes committed and pushed to your copied repository for grading and evaluation. For the assignments, you may not be able to complete all tasks and have all of the tests successfully finishing. This is ok. However, you should endeavor to have as many of the tasks completed before the deadline for the assignment as possible. Also, try and make sure that you only push commits that are building and able to run the tests. You may loose points for pushing a broken build, especially if the last build you submit is not properly compiling and running the tests.

In this problem, up to 25 points will be given for having at least 1 commit that compiles and runs the tests (and at least some attempt was made to work on the first task). Thereafter 5 to 10 points are awarded for completing each of the remaining 6tasks. However you should note that the autograder awards either all point for passing all tests, or no points if any test is failing for one of the tasks. Also note that even if you pass all tests, when the instructor evaluates your assignment, they may remove points if you don't follow the requirements for implementing the code (e.g. must reuse functions here as described, need to correctly declare parameters or member functions as `const` where needed, must have function documentation correct). You may also loose points for style issues. The instructor may give back comments in pull requests and/or create new issues for you if you have issues such as these, so it is good to have work committed early before the due date, so that the instructor may give feedback requesting you to fix issues with your current submission.

Requirements and Grading Rubrics

Program Execution, Output and Functional Requirements

1. Your program must compile, run and produce some sort of output to be graded. 0 if not satisfied.

2. 40 points for keeping code that compiles and runs. A minimum of 50 points will be given if at least the first task is completed and passing tests.
3. 5 to 10 points are awarded for completing each subsequent task 2-10.
4. +5 bonus pts if all system tests pass and your process simulator produces correct output for the given system tests.

Program Style and Documentation

This section is supplemental for the first assignment. If you use the VS Code editor as described for this class, part of the configuration is to automatically run the `clang-format` code style checker/formatter on your code files every time you save the file. You can run this tool manually from the command line as follows:

```
$ make format
clang-format -i include/*.hpp src/*.cpp
```

Class style guidelines have been defined for this class. The `uncrustify.cfg` file defines a particular code style, like indentation, where to place opening and closing braces, whitespace around operators, etc. By running the beautifier on your files it reformats your code to conform to the defined class style guidelines. The beautifier may not be able to fix all style issues, so I might give comments to you about style issues to fix after looking at your code. But you should pay attention to the formatting of the code style defined by this configuration file.

Another required element for class style is that code must be properly documented. Most importantly, all functions and class member functions must have function documentation preceding the function. These have been given to you for the first assignment, but you may need to provide these for future assignment. For example, the code documentation block for the first function you write for this assignment looks like this:

```
/**
 * @brief initialize memory
 *
 * Initialize the contents of memory. Allocate array large enough to
 * hold memory contents for the program. Record base and bounds
 * address for memory address translation. This memory function
 * dynamically allocates enough memory to hold the addresses for the
 * indicated begin and end memory ranges.
 *
 * @param memoryBaseAddress The int value for the base or beginning
 * address of the simulated memory address space for this
 * simulation.
 * @param memoryBoundsAddress The int value for the bounding address,
 * e.g. the maximum or upper valid address of the simulated memory
 * address space for this simulation.
 *
 * @exception Throws SimulatorException if
 * address space is invalid. Currently we support only 4 digit
 * opcodes YYYY, where the 3 digit YYY specifies a reference
 * address. Thus we can only address memory from 000 - 999
 * given the limits of the expected opcode format.
 */
```

This is an example of a `doxygen` formatted code documentation comment. The two `**` starting the block comment are required for `doxygen` to recognize this as a documentation comment. The `@brief`, `@param`, `@exception` etc. tags are used by `doxygen` to build reference documentation from your code. You can build the documentation using the `make docs` build target, though it does require you to have `doxygen` tools installed on your system to work.

```
$ make reldocs
Generating doxygen documentation...
doxygen config/Doxyfile 2>&1 | grep -A 1 warning | egrep -v "assg.*\.md" | grep -v "Found unknown command"
Doxygen version used: 1.9.1
```

The result of this is two new subdirectories in your current directory named `html` and `latex`. You can use a regular

browser to browse the html based documentation in the `html` directory. You will need `latex` tools installed to build the `pdf` reference manual in the `latex` directory.

You can use the `make refdocs` to see if you are missing any required function documentation or tags in your documentation. For example, if you remove one of the `@param` tags from the above function documentation, and run the docs, you would see

```
$ make refdocs
doxygen config/Doxyfile 2>&1 | grep -A 1 warning | egrep -v "assg.*\.md" | grep -v "Found unknown command"

HypotheticalMachineSimulator.hpp:88: warning: The following parameter of
HypotheticalMachineSimulator::initializeMemory(int memoryBaseAddress,
    int memoryBoundsAddress) is not documented:
    parameter 'memoryBoundsAddress'
```

The documentation generator expects that there is a description, and that all input parameters and return values are documented for all functions, among other things. You can run the documentation generation to see if you are missing any required documentation in you project files.